

podcast3

May 9, 2021

```
In [177]: import os
          from scipy.io import wavfile
          import pandas as pd
          import matplotlib.pyplot as plt
          import numpy as np
          from keras.layers import Conv2D, MaxPool2D, Flatten
          from keras.layers import Dropout, Dense, TimeDistributed
          from keras.models import Sequential
          from keras.utils import to_categorical
          from sklearn.utils.class_weight import compute_class_weight
          from tqdm import tqdm
          from python_speech_features import mfcc
          from sklearn.metrics import accuracy_score
          from keras.callbacks import ModelCheckpoint

          path = 'C://Users//richard//OneDrive//Documents//Audacity//wavfiles//'
          audio_classes_path = 'C://Users//richard//OneDrive//Documents//Audacity//audio_classes'
```

0.1 Method

- Using the csv file which contains the list of files and their labels, generate a pandas dataframe.
- Use the dataframe to index the files and get a random window of audio
- Train on this audio using a CNN. Train, test split is not required as all the audio comes from the same place and is of small size
- Go through all the files every 0.1 s and compare prediction to label
- See Final preedcitons and conclusions at bottom of file.

0.2 Class distribution

```
In [178]: df = pd.read_csv(audio_classes_path)
          df.set_index('file_name', inplace=True)

          for f in df.index:
              rate, signal = wavfile.read(path+f)
              df.at[f, 'length'] = signal.shape[0]/rate
```


```

classes = list(np.unique(df.label))
class_dist = df.groupby(['label'])['length'].mean()

fig, ax = plt.subplots()
ax.set_title('Class Distribution', y=1.08)
ax.pie(class_dist, labels=class_dist.index, autopct='%1.1f%%',
        shadow=False, startangle=90)
ax.axis('equal')
plt.show()

```

C:\Users\richard\Anaconda3\envs\tf15\lib\site-packages\scipy\io\wavfile.py:273: WavFileWarning
WavFileWarning)



In [179]: df.head()

Out[179]:

file_name	label	length
music_plus_noise1.wav	music_noise	7.543288
music_plus_noise2.wav	music_noise	6.991315
music_twintones1.wav	music_noise	5.944308
music_plus_speech1.wav	music_speech	4.768209
music1.wav	music	2.018050

In [180]: *#get number of samples to draw from the audio files*
n_samples = 2 * int(df['length'].sum()/0.1)
n_samples

```
Out[180]: 1896
```

```
In [181]: #total length of all the audio in the files
df.length.sum()
```

```
Out[181]: 94.87106575963719
```

```
In [182]: #0 -1 class dist
prob_dist = class_dist / class_dist.sum()

#get a choice of instrument randomly based on class distribution
choices = np.random.choice(class_dist.index,p=prob_dist)

class_dist=class_dist.sort_values(ascending=False)
class_dist
```

```
Out[182]: label
music_noise      6.826304
music_speech     4.768209
speech_nb_noise  2.947545
noise            2.823212
speech_wb_noise  2.554422
music            1.955057
speech_nb        1.409063
speech_wb        1.223050
Name: length, dtype: float64
```

```
In [183]: class_dist.index
```

```
Out[183]: Index(['music_noise', 'music_speech', 'speech_nb_noise', 'noise',
                'speech_wb_noise', 'music', 'speech_nb', 'speech_wb'],
                dtype='object', name='label')
```

0.3 Make a CNN model

Use general image processing network, increasing size of feature maps at each layer. As output is eight classes use softmax at output to get majority class probability. Max pooling can be used but images are small anyway 13x13 or 26x26. Since using small size and same data files will overfit anyway so not using test and validation set.

```
In [184]: def get_conv_model():
    model = Sequential()
    model.add(Conv2D(32,(3,3),activation='relu',strides=(1,1),padding='same',input_shape=(1,13,13)))
    #model.add(MaxPool2D((2,2)))
    model.add(Conv2D(64,(3,3),activation='relu',strides=(1,1),padding='same'))
    #model.add(MaxPool2D((2,2)))
    model.add(Conv2D(128,(3,3),activation='relu',strides=(1,1),padding='same'))
    #model.add(MaxPool2D((2,2)))
    model.add(Conv2D(256,(3,3),activation='relu',strides=(1,1),padding='same'))
```

```

#model.add(MaxPool2D((2,2)))
model.add(Dropout(0.1))
model.add(Flatten())
model.add(Dense(256,activation='relu'))
model.add(Dense(128,activation='relu'))
model.add(Dense(64,activation='relu'))
model.add(Dense(8,activation='softmax'))
model.summary()
model.compile(loss='categorical_crossentropy',optimizer='adam',metrics=['acc'])
return model

```

0.4 Make a helper class

```

In [185]: class Config():
    def __init__(self,mode = 'conv', nfilt=26, nfeat=26, nfft=2048,rate=16000):
        self.mode = mode
        self.nfilt = nfilt
        self.nfeat = nfeat
        self.nfft = nfft
        self.rate = rate
        self.step = int(rate/10)

```

0.5 Generate data from files

Need to get random chunks of the data from the files. These chunks here are 0.1s long and selected randomly from the distribution given above

```

In [186]: #generate random sampling from dataset
def build_rand_feat():
    X = [];y=[]

    _min,_max = float('inf'),-float('inf')

    for _ in tqdm(range(n_samples)):
        rand_class = np.random.choice(class_dist.index,p=prob_dist)

        #get random file based on random instrument
        file = np.random.choice(df[df.label==rand_class].index)
        rate,wav = wavfile.read(path + file)
        label = df.at[file,'label']

        #get random chunk from file
        rand_index = np.random.randint(0,wav.shape[0] - config.step)
        sample = wav[rand_index:rand_index+config.step]
        X_sample = mfcc(sample,rate,numcep=config.nfeat, nfilt = config.nfilt, nfft=

        #running update of min/max
        _min = min(np.amin(X_sample),_min)

```

```

        _max = max(np.amax(X_sample), _max)
        X.append(X_sample)
        y.append(classes.index(label)) #get index number

    config.min = _min
    config.max = _max

    X, y = np.array(X), np.array(y) #normalise the data

    X = (X - _min) / (_max - _min)

    X = X.reshape(X.shape[0], X.shape[1], X.shape[2], 1) #samples * rows (time) * cols
    y = to_categorical(y, num_classes=8)

    _min = str(_min)
    _max = str(_max)

    return X, y, _min, _max

```

In [187]: config= Config()

In [188]: *#build dataset using distribution we have defined*

```

X, y, _min, _max = build_rand_feat()
y_flat = np.argmax(y, axis=1) #get the max value across each row
input_shape = (X.shape[1], X.shape[2], 1)
model = get_conv_model()

```

3%| | 63/1896 [00:05<00:00, 317.53it/s]

WavFileWarning)

100%| 1896/1896 [00:05<00:00, 317.53it/s]

Layer (type)	Output Shape	Param #
conv2d_29 (Conv2D)	(None, 3, 26, 32)	320
conv2d_30 (Conv2D)	(None, 3, 26, 64)	18496
conv2d_31 (Conv2D)	(None, 3, 26, 128)	73856
conv2d_32 (Conv2D)	(None, 3, 26, 256)	295168
dropout_8 (Dropout)	(None, 3, 26, 256)	0
flatten_8 (Flatten)	(None, 19968)	0
dense_29 (Dense)	(None, 256)	5112064

dense_30 (Dense)	(None, 128)	32896

dense_31 (Dense)	(None, 64)	8256

dense_32 (Dense)	(None, 8)	520
=====		
Total params: 5,541,576		
Trainable params: 5,541,576		
Non-trainable params: 0		

```
In [189]: print(X.shape)
          print(y.shape)
```

```
(1896, 3, 26, 1)
(1896, 8)
```

```
In [190]: #base weights on prob dist defined above - need as unbalanced
          class_weight = compute_class_weight('balanced',np.unique(y_flat),y_flat)
```

```
In [191]: class_weight
```

```
Out[191]: array([1.05803571, 1.52903226, 0.43807763, 1.15048544, 2.69318182,
                 0.62204724, 1.16748768, 2.41836735])
```

```
In [192]: #try adding weights to classes or not
          #model.fit(X,y,epochs=30,batch_size=32,shuffle=True,class_weight=class_weight,verbose=2)
          model.fit(X,y,epochs=50,batch_size=16,shuffle=True,verbose=2)
```

```
Epoch 1/50
- 2s - loss: 1.8995 - acc: 0.2911
Epoch 2/50
- 2s - loss: 0.7511 - acc: 0.7431
Epoch 3/50
- 2s - loss: 0.5452 - acc: 0.8043
Epoch 4/50
- 2s - loss: 0.4738 - acc: 0.8233
Epoch 5/50
- 2s - loss: 0.3807 - acc: 0.8523
Epoch 6/50
- 2s - loss: 0.3445 - acc: 0.8745
Epoch 7/50
- 2s - loss: 0.3194 - acc: 0.8739
Epoch 8/50
- 2s - loss: 0.3054 - acc: 0.8797
Epoch 9/50
- 2s - loss: 0.2588 - acc: 0.9024
```

Epoch 10/50
- 2s - loss: 0.2461 - acc: 0.9103
Epoch 11/50
- 2s - loss: 0.2229 - acc: 0.9182
Epoch 12/50
- 2s - loss: 0.2072 - acc: 0.9272
Epoch 13/50
- 2s - loss: 0.2244 - acc: 0.9204
Epoch 14/50
- 2s - loss: 0.1882 - acc: 0.9330
Epoch 15/50
- 2s - loss: 0.1687 - acc: 0.9346
Epoch 16/50
- 2s - loss: 0.1273 - acc: 0.9525
Epoch 17/50
- 2s - loss: 0.1390 - acc: 0.9499
Epoch 18/50
- 2s - loss: 0.1422 - acc: 0.9509
Epoch 19/50
- 2s - loss: 0.1192 - acc: 0.9578
Epoch 20/50
- 2s - loss: 0.1400 - acc: 0.9515
Epoch 21/50
- 2s - loss: 0.1037 - acc: 0.9641
Epoch 22/50
- 2s - loss: 0.0976 - acc: 0.9673
Epoch 23/50
- 2s - loss: 0.1026 - acc: 0.9673
Epoch 24/50
- 2s - loss: 0.1040 - acc: 0.9631
Epoch 25/50
- 2s - loss: 0.1080 - acc: 0.9620
Epoch 26/50
- 2s - loss: 0.0553 - acc: 0.9815
Epoch 27/50
- 2s - loss: 0.0515 - acc: 0.9826
Epoch 28/50
- 2s - loss: 0.0357 - acc: 0.9852
Epoch 29/50
- 2s - loss: 0.0706 - acc: 0.9784
Epoch 30/50
- 2s - loss: 0.1031 - acc: 0.9626
Epoch 31/50
- 2s - loss: 0.0647 - acc: 0.9768
Epoch 32/50
- 2s - loss: 0.0695 - acc: 0.9747
Epoch 33/50
- 2s - loss: 0.0874 - acc: 0.9678

```

Epoch 34/50
- 2s - loss: 0.0649 - acc: 0.9800
Epoch 35/50
- 2s - loss: 0.0320 - acc: 0.9884
Epoch 36/50
- 2s - loss: 0.0560 - acc: 0.9836
Epoch 37/50
- 2s - loss: 0.0389 - acc: 0.9868
Epoch 38/50
- 2s - loss: 0.0711 - acc: 0.9773
Epoch 39/50
- 2s - loss: 0.0699 - acc: 0.9757
Epoch 40/50
- 2s - loss: 0.0243 - acc: 0.9916
Epoch 41/50
- 2s - loss: 0.0164 - acc: 0.9931
Epoch 42/50
- 2s - loss: 0.0265 - acc: 0.9910
Epoch 43/50
- 2s - loss: 0.0531 - acc: 0.9805
Epoch 44/50
- 2s - loss: 0.0496 - acc: 0.9826
Epoch 45/50
- 2s - loss: 0.1470 - acc: 0.9573
Epoch 46/50
- 2s - loss: 0.0508 - acc: 0.9831
Epoch 47/50
- 2s - loss: 0.0186 - acc: 0.9942
Epoch 48/50
- 2s - loss: 0.0702 - acc: 0.9763
Epoch 49/50
- 2s - loss: 0.0407 - acc: 0.9873
Epoch 50/50
- 2s - loss: 0.0385 - acc: 0.9900

```

```
Out[192]: <keras.callbacks.History at 0x18d95d8a908>
```

1 Predictions 1

Use the model to get predictions, first across all files in directory

```

In [193]: def build_pred_all_files(fn2class,audio_dir=path):
            y_true = []
            y_pred = []
            fn_prob = {}

            for fn in tqdm(os.listdir(audio_dir)):

```



```

rate, wav = wavfile.read(os.path.join(audio_dir,fn))
label = fn2class[fn]
c = classes.index(label)
y_prob = []
y_per_file=[]

#print(wav.shape[0],config.step,fn)
#break

for i in range(0,wav.shape[0]-config.step,config.step):
    sample = wav[i:i+config.step]
    #sample is a block of samples x secs (0.1) long
    x = mfcc(sample,rate, numcep=config.nfeat,nfilt = config.nfilt,nfft=config.nfft)
    x = (x-float(_min))/(float(_max)-float(_min))
    x = x.reshape(1,x.shape[0],x.shape[1],1)

    #for every block of samples get the mfcc and compare to target
    #y_hat is probability for each class - softmax in final layer
    y_hat = model.predict(x)
    y_prob.append(y_hat)
    #y_pred is class corresponding to max prob. for each block of samples (even if not a class)
    y_pred.append(np.argmax(y_hat))
    #y_true is the real class
    y_true.append(c)

#print(label,c)
fn_prob[fn] = np.mean(y_prob,axis=0).flatten()

return y_true, y_pred, fn_prob

```

In [194]: *#y_pred is the result for each block of samples for every file*
#y_true is the actual value for each block of samples for every file
#fn_prob is the value for each class probability for each file, the max value indicates the predicted class

```

df = pd.read_csv(audio_classes_path)

classes = list(np.unique(df.label))

fn2class = dict(zip(df.file_name,df.label))

y_true, y_pred, fn_prob = build_pred_all_files(fn2class,path)
result = accuracy_score(y_true,y_pred)

```

```

22%|          | 8/37 [00:03<00:13,  2.1it/s]
WavFileWarning)
84%|          | 31/37 [00:09<00:01,  3.11it/s]C:\Users\richard\Anaconda3\envs\tf15\lib\site-packages\tqdm\tqdm.py:101:
TqdmSynchronisationWarning)
100%|| 37/37 [00:11<00:00,  3.12it/s]

```

```
In [195]: #go through every file point x seconds and get total accuracy overall
print("The overall accuracy every 0.1 secs across all wav files is {:.2f}".format(res))
```

The overall accuracy every 0.1 secs across all wav files is 0.82

1.1 Prediction 2

Label each file using the majority class per all 0.1 seconds per file. Fo example is music1.wav labelled as music?

```
In [205]: #write result for each file usually 100% correct?
if write_to_output:
    y_probs = []
    for i, row in df.iterrows():
        y_prob = fn_prob[row.file_name]
        y_probs.append(y_prob)
        for c,p in zip(classes,y_prob):
            df.at[i,c] = p
    y_pred = [ classes [ np.argmax(y)] for y in y_probs]
    df['y_pred'] = y_pred
    #get rid of the probabilities for write to file
    df = df[["file_name", "label",'y_pred']]
    #df.to_csv('predictions.csv',index=False)
df
```

```
Out[205]:
```

	file_name	label	y_pred
0	music_plus_noise1.wav	music_noise	music_noise
1	music_plus_noise2.wav	music_noise	music_noise
2	music_twintones1.wav	music_noise	music_noise
3	music_plus_speech1.wav	music_speech	music_speech
4	music1.wav	music	music
5	music2.wav	music	music
6	music3.wav	music	music
7	music4.wav	music	music
8	music5.wav	music	music
9	noise1.wav	noise	noise
10	noise2.wav	noise	noise
11	noise3.wav	noise	noise
12	noise4.wav	noise	noise
13	noise5.wav	noise	noise
14	noise6.wav	noise	noise
15	speech_nb_plus_noise1.wav	speech_nb_noise	speech_nb_noise
16	speech_nb_plus_noise2.wav	speech_nb_noise	speech_nb_noise
17	speech_nb_plus_noise3.wav	speech_nb_noise	speech_nb_noise
18	speech_nb_plus_noise4.wav	speech_nb_noise	speech_nb_noise
19	speech_nb1.wav	speech_nb	speech_nb
20	speech_nb2.wav	speech_nb	speech_nb
21	speech_nb3.wav	speech_nb	speech_nb

22	speech_nb4.wav	speech_nb	speech_nb
23	speech_nb5.wav	speech_nb	speech_nb
24	speech_nb6.wav	speech_nb	speech_nb
25	speech_wb_plus_noise1.wav	speech_wb_noise	speech_wb_noise
26	speech_wb_plus_noise2.wav	speech_wb_noise	speech_wb_noise
27	speech_wb_plus_noise3.wav	speech_wb_noise	speech_wb_noise
28	speech_wb_plus_noise4.wav	speech_wb_noise	speech_wb_noise
29	speech_wb_plus_noise5.wav	speech_wb_noise	speech_wb_noise
30	speech_wb_plus_noise6.wav	speech_wb_noise	speech_wb
31	speech_wb1.wav	speech_wb	speech_wb
32	speech_wb2.wav	speech_wb	speech_wb
33	speech_wb3.wav	speech_wb	speech_wb
34	speech_wb4.wav	speech_wb	speech_wb
35	speech_wb5.wav	speech_wb	speech_wb
36	speech_wb6.wav	speech_wb	speech_wb

All files are correctly labelled as the correct type. Doesn't agree with results below!

2 Final Predictions

Get the predictions for every x samples and compare for each file. These are the final and most appropriate results

```
In [197]: def build_pred_per_file(fn2class, audio_dir, fn):
    y_true = []
    y_pred = []

    #for fn in tqdm(os.listdir(audio_dir)):
    rate, wav = wavfile.read(os.path.join(audio_dir, fn))
    label = fn2class[fn]
    c = classes.index(label)

    for i in range(0, wav.shape[0]-config.step, config.step):
        sample = wav[i:i+config.step]
        #sample is a block of samples x secs long
        x = mfcc(sample, rate, numcep=config.nfeat, nfilt = config.nfilt, nfft=config.nfft)
        x = (x - float(_min)) / (float(_max) - float(_min))
        x = x.reshape(1, x.shape[0], x.shape[1], 1)
        #for every block of samples get the mfcc and compare to target
        #y_hat is probability for each class - softmax in final layer
        y_hat = model.predict(x)
        #y_pred is class corresponding to max prob. for each block of samples (every x)
        y_pred.append(np.argmax(y_hat))
        #y_true is the real class
        y_true.append(c)

    return y_true, y_pred
```

```

In [198]: files = os.listdir(path)
          sum_scores = 0
          for file in files:
              y_true, y_pred = build_pred_per_file(fn2class,path,file)
              score = accuracy_score(y_true,y_pred)
              sum_scores += score;
              print("The overall accuracy score per 0.1 secs is {:.2f} for {}".format(score,file))

          print(sum_scores/len(files))

```

```

The overall accuracy score per 0.1 secs is 0.78 for music1.wav
The overall accuracy score per 0.1 secs is 1.00 for music2.wav
The overall accuracy score per 0.1 secs is 0.93 for music3.wav
The overall accuracy score per 0.1 secs is 0.98 for music4.wav
The overall accuracy score per 0.1 secs is 0.67 for music5.wav
The overall accuracy score per 0.1 secs is 0.92 for music_plus_noise1.wav
The overall accuracy score per 0.1 secs is 0.56 for music_plus_noise2.wav
The overall accuracy score per 0.1 secs is 0.98 for music_plus_speech1.wav

```

```

C:\Users\richard\Anaconda3\envs\tf15\lib\site-packages\scipy\io\wavfile.py:273: WavFileWarning
WavFileWarning)

```

```

The overall accuracy score per 0.1 secs is 0.87 for music_twintones1.wav
The overall accuracy score per 0.1 secs is 1.00 for noise1.wav
The overall accuracy score per 0.1 secs is 1.00 for noise2.wav
The overall accuracy score per 0.1 secs is 1.00 for noise3.wav
The overall accuracy score per 0.1 secs is 1.00 for noise4.wav
The overall accuracy score per 0.1 secs is 1.00 for noise5.wav
The overall accuracy score per 0.1 secs is 1.00 for noise6.wav
The overall accuracy score per 0.1 secs is 0.82 for speech_nb1.wav
The overall accuracy score per 0.1 secs is 0.87 for speech_nb2.wav
The overall accuracy score per 0.1 secs is 0.96 for speech_nb3.wav
The overall accuracy score per 0.1 secs is 0.75 for speech_nb4.wav
The overall accuracy score per 0.1 secs is 0.78 for speech_nb5.wav
The overall accuracy score per 0.1 secs is 0.83 for speech_nb6.wav
The overall accuracy score per 0.1 secs is 0.98 for speech_nb_plus_noise1.wav
The overall accuracy score per 0.1 secs is 0.99 for speech_nb_plus_noise2.wav
The overall accuracy score per 0.1 secs is 0.96 for speech_nb_plus_noise3.wav
The overall accuracy score per 0.1 secs is 0.85 for speech_nb_plus_noise4.wav
The overall accuracy score per 0.1 secs is 0.87 for speech_wb1.wav
The overall accuracy score per 0.1 secs is 0.97 for speech_wb2.wav
The overall accuracy score per 0.1 secs is 0.83 for speech_wb3.wav
The overall accuracy score per 0.1 secs is 0.86 for speech_wb4.wav
The overall accuracy score per 0.1 secs is 0.95 for speech_wb5.wav
The overall accuracy score per 0.1 secs is 0.96 for speech_wb6.wav
The overall accuracy score per 0.1 secs is 0.57 for speech_wb_plus_noise1.wav

```

The overall accuracy score per 0.1 secs is 0.62 for speech_wb_plus_noise2.wav
The overall accuracy score per 0.1 secs is 0.62 for speech_wb_plus_noise3.wav
The overall accuracy score per 0.1 secs is 0.54 for speech_wb_plus_noise4.wav
The overall accuracy score per 0.1 secs is 0.64 for speech_wb_plus_noise5.wav
The overall accuracy score per 0.1 secs is 0.13 for speech_wb_plus_noise6.wav
0.8385432753854886

3 Conclusions

- We can see that the results for tend to be much better for single component audio than audio composed of two components. For example noise and speech_wb_plus_noise.
- The cause of the error is the way the files were made by me and the sample window used to collect the audio
- When the files were made I mixed two components together to generate the audio file. For example speech plus noise was constructed using speech and twin tone noise. Therefore in that file speech is continuous but varies in amplitude but the twin tone lasts only for a short time. When I take a random sample of the audio file I may not be taking speech AND twin tone - It may be just speech or speech AND small component of twin tone. I believe this is the main cause of errors you see here.
- This is also seen when making predictions, two component audio results can sometimes be very good - it just depends if the sample collection is capturing the two audio components.